

# ResQ-Graph: Autonomous Fleet Optimization System

## Agile Sprint Schedule (14 Weeks)

**Project Duration:** 14 weeks | **Sprint Length:** 1 week | **Total Sprints:** 14 **Team Load:** 6-course load + project (light to moderate commitment per sprint)

---

## PHASE 1: THE SKELETON (Weeks 1-3)

*Goal: Establish foundation with working map and basic navigation*

### Sprint 1: Project Setup & Map Infrastructure (Week 1)

**Sprint Goal:** Set up development environment and prepare baseline map data

#### User Stories:

##### 1. US-001: Initialize Development Environment

- As a developer, I want all required libraries installed and configured so that I can begin coding
- Acceptance Criteria:
  - pip environment set up with osmnx, networkx, numpy, matplotlib
  - All imports tested and verified working
  - Project structure created (src/, tests/, data/, docs/ folders)
- Estimated Points: 3
- Tasks: Install packages, create folder structure, verify imports

##### 2. US-002: Select and Validate Sandbox Map

- As a project lead, I want a well-defined, manageable sandbox area so that testing is efficient
- Acceptance Criteria:
  - Model Town, Lahore selected as primary sandbox
  - Map boundaries documented (lat/lon coordinates)
  - OSM data quality verified (no isolated islands expected)
- Estimated Points: 2
- Tasks: Research Model Town, identify boundaries, validate OSM coverage

##### 3. US-003: Implement Map Baking Pipeline (bake\_map.py)

- As a system architect, I want raw OSM data converted to a clean graph so that algorithms can operate efficiently

- Acceptance Criteria:
  - OSM data downloaded for Model Town
  - Isolated nodes removed (connectivity check enforced)
  - Graph exported as modeltown.graphml
  - Node/edge count and graph statistics logged
- Estimated Points: 5
- Tasks: Write OSMnx query, implement isolation removal, save graphml format

#### 4. US-004: Create Distance Matrix Utility Module

- As a developer, I want a reusable function to compute and cache distances so that lookups are O(1)
  - Acceptance Criteria:
    - Floyd-Warshall or NetworkX shortest path pre-computed
    - Distance matrix saved as .npy file
    - Loading/caching function created
    - Verification script confirms matrix correctness
  - Estimated Points: 5
  - Tasks: Implement matrix computation, create serialization, write verification tests
- 

### Sprint 2: Graph Navigation & A\* Search (Week 2)

**Sprint Goal:** Implement custom A\* pathfinder and validate on real map

#### User Stories:

1. *US-005: Implement Custom A Pathfinding Algorithm\**
  - As an AI engineer, I want A\* search implemented from scratch so that I understand the algorithm and can extend it
  - Acceptance Criteria:
    - Priority queue-based A\* written without nx.shortest\_path
    - Haversine heuristic implemented (lat/lon coordinates)
    - Function returns ordered list of nodes from start to goal
    - Handles edge cases: unreachable nodes, same start/goal, empty graph
  - Estimated Points: 8
  - Tasks: Write priority queue logic, implement heuristic, handle edge cases, code review
2. *US-006: Validate A with Manual Test Cases\**

- As a QA engineer, I want concrete test cases so that A\* correctness is verified
- Acceptance Criteria:
  - 5 test cases defined: short path, long path, unreachable, same node, complex routing
  - Visual validation: plot start, goal, and path on map
  - Path validity confirmed (all edges exist in graph)
  - Performance logged: path computation < 50ms on sandbox
- Estimated Points: 4
- Tasks: Define test cases, write validation script, visualize paths, benchmark

### 3. US-007: Create Path Visualization Module

- As a developer, I want to visualize paths on the map so that correctness is obvious
- Acceptance Criteria:
  - Matplotlib-based visualization created
  - Start node marked (green), goal marked (red), path in blue
  - Map background shows street network
  - Interactivity: zoom/pan functional
- Estimated Points: 4
- Tasks: Write plotting function, style map layers, test interactivity

### 4. US-008: Document Navigation Layer API

- As a future maintainer, I want clear API documentation so that the module is reusable
  - Acceptance Criteria:
    - Function signatures documented (docstrings)
    - Algorithm explanation (pseudocode)
    - Example usage script provided
    - README updated with navigation module overview
  - Estimated Points: 2
  - Tasks: Write docstrings, create examples, update README
- 

## Sprint 3: Genetic Algorithm & Strategic Solver (Week 3)

**Sprint Goal:** Implement facility location optimization using evolutionary algorithm

**User Stories:**

### 1. US-009: Implement Genetic Algorithm Framework

- As an optimization engineer, I want a generic GA so that base station locations are evolved
- Acceptance Criteria:
  - Genome representation: list of 5 node IDs
  - Population initialization: 50 random genomes
  - Selection mechanism: fitness-proportionate (roulette wheel)
  - Mutation: random node swap
  - Crossover: single-point crossover
- Estimated Points: 8
- Tasks: Implement GA class, write fitness evaluation, implement operators

## 2. US-010: Implement Fitness Function (Facility Location)

- As a data scientist, I want a fitness function that minimizes total travel time so that GA evolves good solutions
- Acceptance Criteria:
  - Fitness = sum of distances from all nodes to nearest ambulance base
  - Distances pre-computed from distance matrix ( $O(1)$  lookup)
  - Handles edge cases: duplicate bases penalized, all nodes reachable
  - Fitness values logged per generation
- Estimated Points: 5
- Tasks: Implement fitness calculation, add penalty logic, test with known solutions

## 3. US-011: Run GA to Convergence & Track Progress

- As a solver developer, I want GA to converge and visualize evolution so that solution quality is observable
- Acceptance Criteria:
  - GA runs for 100 generations
  - Best fitness logged per generation
  - Convergence plot created (matplotlib)
  - Final 5 station locations recorded and exported
  - Execution time  $< 2$  minutes on sandbox
- Estimated Points: 4
- Tasks: Write GA execution loop, create convergence tracker, plot results

## 4. US-012: Visualize Optimal Station Locations on Map

- As a product manager, I want to see where stations are placed so that the solution is interpretable

- Acceptance Criteria:
    - Map plotted with street network
    - Optimal stations marked as blue stars
    - Coverage zones visualized (Voronoi or heatmap)
    - Comparison: random stations (gray) vs. optimal (blue)
    - HTML report generated with visualization
  - Estimated Points: 4
  - Tasks: Implement Voronoi/coverage visualization, create comparison plot, generate report
- 

## PHASE 2: THE BRAIN (Weeks 4-7)

*Goal: Add simulation engine and dynamic intelligence*

### Sprint 4: Simulation Engine & Ambulance Agents (Week 4)

**Sprint Goal:** Create basic simulation loop with mobile ambulance agents

#### User Stories:

##### 1. US-013: Implement Ambulance Agent Class

- As a simulation developer, I want an Ambulance class so that individual agents are represented
- Acceptance Criteria:
  - State machine: IDLE, IN\_TRANSIT, ON\_SCENE
  - Properties: current\_location, current\_path, assigned\_task
  - Methods: navigate(destination), update\_position(), get\_status()
  - Location updated based on A\* path each simulation tick
- Estimated Points: 5
- Tasks: Design state machine, implement class, write unit tests

##### 2. US-014: Implement Event Spawner (Poisson Distribution)

- As a simulation engineer, I want random accident generation so that emergency calls are realistic
- Acceptance Criteria:
  - Poisson process generates accidents with configurable  $\lambda$  (rate parameter)
  - Accidents spawn at random nodes
  - Accident attributes: timestamp, location, priority
  - Spawning logic isolated in separate function for easy modification

- Estimated Points: 4
- Tasks: Implement Poisson logic, test distribution, create configurability

### 3. US-015: Build Main Simulation Loop

- As a simulator, I want a tick-based engine so that time-stepped physics works
- Acceptance Criteria:
  - Main loop processes N ticks (configurable)
  - Each tick: (1) update ambulance positions, (2) spawn events, (3) assign idle ambulances
  - Simulation state tracked (current\_tick, active\_events, ambulance\_positions)
  - Loop runs without errors for 1000+ ticks
- Estimated Points: 6
- Tasks: Design loop architecture, implement tick logic, handle state management

### 4. US-016: Create Basic Live Visualization (Matplotlib Interactive)

- As a developer, I want to see the simulation running so that behavior is observable
  - Acceptance Criteria:
    - Interactive matplotlib plot created (plt.ion() mode)
    - Ambulances shown as colored dots on map
    - Accidents shown as red X markers
    - Plot redraws every N ticks (configurable)
    - FPS sufficient for real-time observation
  - Estimated Points: 4
  - Tasks: Implement interactive plotting, optimize refresh rate, test stability
- 

## Sprint 5: Dispatcher Logic & Task Assignment (Week 5)

**Sprint Goal:** Implement dispatch brain that assigns ambulances to emergencies

**User Stories:**

### 1. US-017: Implement Dispatcher Brain Class

- As a system architect, I want a centralized dispatcher so that assignment decisions are coordinated
- Acceptance Criteria:
  - DispatcherBrain class created
  - Methods: assign\_task(ambulance, event), rebalance\_fleet()
  - Tracks active events and idle ambulances

- Decision logic: assign nearest idle ambulance to new event
- Estimated Points: 5
- Tasks: Design dispatcher class, implement assignment logic, write tests

## 2. US-018: Implement Task Assignment Algorithm

- As an operations researcher, I want optimal assignment so that response times are minimized
- Acceptance Criteria:
  - Algorithm: find nearest idle ambulance to event using distance matrix
  - Ties broken by ambulance\_id for determinism
  - Assignment creates navigation path using A\*
  - Event marked as "assigned" in dispatcher state
- Estimated Points: 4
- Tasks: Implement assignment, integrate with A\*, test edge cases

## 3. US-019: Track Response Time Metrics

- As a data analyst, I want metrics so that performance is measurable
- Acceptance Criteria:
  - Response time = spawn\_time - arrival\_time
  - All response times logged per event
  - Average Response Time (ART) computed after simulation
  - Metrics exported to CSV for analysis
- Estimated Points: 3
- Tasks: Implement metrics tracking, add logging, export function

## 4. US-020: Integrate Dispatcher into Simulation Loop

- As an engineer, I want dispatcher integrated so that the full system works
- Acceptance Criteria:
  - Dispatcher called each tick
  - New events auto-assigned to idle ambulances
  - Ambulances navigate using A\* paths
  - Simulation runs with full dispatch logic for 1000 ticks
- Estimated Points: 3
- Tasks: Integrate components, test full loop, debug interaction

---

## Sprint 6: K-Means Clustering & Dynamic Hotspot Detection (Week 6)

**Sprint Goal:** Implement unsupervised learning for demand prediction

### User Stories:

#### 1. US-021: Implement K-Means Clustering from Scratch

- As a data scientist, I want K-Means implemented manually so that I understand the algorithm
- Acceptance Criteria:
  - Manual K-Means written (no sklearn)
  - Centroid initialization: random or k-means++
  - Iterative update: reassign points, recompute centroids
  - Convergence check: centroid movement < epsilon
  - Returns cluster assignments and centroids
- Estimated Points: 8
- Tasks: Implement clustering algorithm, write convergence logic, test on sample data

#### 2. US-022: Create Demand Clustering Module

- As a tactical system, I want to detect accident hotspots so that ambulances pre-position
- Acceptance Criteria:
  - Input: list of active accident locations
  - Process: run K-Means with k=2-3 (configurable)
  - Output: cluster centroids (hotspot locations)
  - Minimum cluster size enforced (to avoid noise)
- Estimated Points: 4
- Tasks: Wrap K-Means for demand data, add filtering, write validation

#### 3. US-023: Integrate Dynamic Hotspot Rebalancing into Dispatcher

- As a dispatcher, I want idle ambulances repositioned so that coverage is reactive
- Acceptance Criteria:
  - Every 50 ticks, run K-Means on active events
  - If hotspots found, move idle ambulances to centroids
  - Ambulances not disrupted if already assigned
  - Rebalancing logged for analysis
- Estimated Points: 5
- Tasks: Create rebalancing function, integrate into dispatcher tick, add logging

#### 4. US-024: Visualize Hotspots on Live Map

- As an observer, I want to see clusters so that behavior is intuitive
  - Acceptance Criteria:
    - Cluster centroids shown as blue circles on map
    - Cluster membership shown (e.g., shaded regions)
    - Hotspot evolution visible over time
    - Overlay on existing ambulance visualization
  - Estimated Points: 4
  - Tasks: Implement cluster visualization, optimize rendering, test with real runs
- 

### Sprint 7: Simulation Refinement & Traffic Dynamics (Week 7)

**Sprint Goal:** Add realism with traffic simulation and edge weight dynamics

#### User Stories:

##### 1. US-025: Implement Traffic Congestion Simulation

- As a realism engineer, I want traffic to slow ambulances so that simulation is realistic
- Acceptance Criteria:
  - Traffic weight multiplier: 1.0 (free flow) to 2.5 (congestion)
  - Congestion zones: high-accident areas incur 2x time penalty
  - Dynamic: congestion increases with event density, decays over time
  - A\* re-routes around congestion (uses current edge weights)
- Estimated Points: 6
- Tasks: Implement congestion model, integrate with graph weights, test rerouting

##### 2. US-026: Implement Ambulance Re-routing on Path

- As a navigator, I want ambulances to re-route around obstacles so that they adapt
- Acceptance Criteria:
  - While ambulance is in transit, check for congestion updates
  - If blocked node's weight increased significantly, recompute path
  - New path computed using A\* and current weights
  - Re-routing logged with reason (e.g., "congestion detected")
- Estimated Points: 4
- Tasks: Implement re-routing check, compute new paths, track re-routes

### 3. US-027: Add Simulation Parameters Configuration File

- As a researcher, I want configurable parameters so that I can run experiments
- Acceptance Criteria:
  - YAML/JSON config file with all parameters ( $\lambda$ , traffic\_factor, k\_update\_interval, etc.)
  - Config loader: parse file and set simulation parameters
  - Documentation: each parameter explained with default/range
  - Example configs: baseline, high-traffic, cluster-heavy
- Estimated Points: 3
- Tasks: Design config schema, implement loader, create example configs

### 4. US-028: Create Comprehensive Simulation Logging System

- As a debugger, I want detailed logs so that I can analyze and debug
- Acceptance Criteria:
  - Event-level logging: spawn, assignment, arrival, completion
  - Ambulance-level logging: state changes, path updates, re-routes
  - Tick-level logging: number of active events, idle ambulances, hotspots
  - Logs written to CSV and live console output
- Estimated Points: 3
- Tasks: Implement logging framework, add log points throughout code, test output

---

## PHASE 3: THE PULSE (Weeks 8-12)

*Goal: Complete simulation, optimize, and run experiments*

### Sprint 8: Baseline Comparison & Random Fleet (Week 8)

**Sprint Goal:** Establish performance baseline with random station placement

#### User Stories:

##### 1. US-029: Implement Random Station Placement Generator

- As a baseline engineer, I want a random fleet so that I have a control group
- Acceptance Criteria:
  - Generate 5 random node IDs from graph (no duplicates)
  - Repeat generation N times (e.g., 10) to get average baseline
  - Random placements logged for reproducibility

- Function returns list of node IDs
- Estimated Points: 2
- Tasks: Implement random selection, test reproducibility

## 2. US-030: Run Simulation with Random Fleet (Baseline)

- As a researcher, I want a baseline ART so that optimization gains are measurable
- Acceptance Criteria:
  - Run simulation 10 times with different random seeds
  - Each run: 1000+ ticks with Poisson event spawning
  - Collect ART (Average Response Time) for each run
  - Calculate mean ART and standard deviation
  - Results exported to CSV
- Estimated Points: 5
- Tasks: Create baseline runner script, run experiments, analyze results

## 3. US-031: Document Random Fleet Baseline Results

- As a analyst, I want clear baseline documentation so that comparisons are meaningful
- Acceptance Criteria:
  - Report: baseline ART, std dev, number of events processed
  - Visualizations: ART distribution (histogram), time-series response times
  - Analysis: observations about random placement inefficiencies
  - Markdown report generated automatically
- Estimated Points: 3
- Tasks: Create analysis script, generate visualizations, write report

## 4. US-032: Create Baseline Configuration & Reproducibility Seed

- As a researcher, I want reproducible results so that experiments are valid
- Acceptance Criteria:
  - Baseline seed documented (random\_seed for random fleet generation)
  - Simulation seeds set (event spawning, ambulance initialization)
  - Config version tracked
  - README documents how to reproduce baseline run
- Estimated Points: 2
- Tasks: Set up seed management, document process

## Sprint 9: AI-Optimized Fleet & Comparison (Week 9)

**Sprint Goal:** Run optimized simulation and compare against baseline

### User Stories:

#### 1. US-033: Run Simulation with AI-Optimized Fleet

- As a researcher, I want AI fleet performance so that optimization impact is measured
- Acceptance Criteria:
  - Load GA-optimal stations from Week 3 results
  - Run simulation 10 times (same seed series as baseline for fair comparison)
  - Each run: 1000+ ticks with identical event distribution
  - Collect ART for each run
  - Results exported to CSV
- Estimated Points: 4
- Tasks: Load optimal stations, run experiment loop, export results

#### 2. US-034: Implement Head-to-Head Comparison Analysis

- As a data scientist, I want statistical comparison so that the improvement is quantified
- Acceptance Criteria:
  - Compare baseline ART vs. AI ART (paired samples)
  - Calculate: absolute improvement, percentage improvement, standard error
  - Statistical test: t-test for significance (p-value)
  - All metrics documented
- Estimated Points: 4
- Tasks: Implement comparison logic, run statistical tests, document results

#### 3. US-035: Generate Comparison Visualizations

- As a presenter, I want visual comparisons so that results are compelling
- Acceptance Criteria:
  - Plot 1: Side-by-side ART distributions (baseline vs. AI)
  - Plot 2: Time-series response times for both fleets
  - Plot 3: Heatmaps showing demand coverage (baseline vs. AI)
  - Plot 4: Station placement comparison (random vs. optimal)
  - All plots with error bars and significance markers
- Estimated Points: 5

- Tasks: Create comparison plots, style for presentation, verify correctness

#### 4. US-036: Create Results Summary Document

- As a reporter, I want a concise summary so that key findings are clear
  - Acceptance Criteria:
    - Executive summary: key metrics and findings
    - Detailed results table: all ART values, improvements
    - Methodology: how experiments were conducted
    - Caveats & limitations noted
    - PDF/Markdown report generated
  - Estimated Points: 3
  - Tasks: Write summary, compile all results, generate PDF
- 

### Sprint 10: Sensitivity Analysis & Parameter Tuning (Week 10)

**Sprint Goal:** Understand how parameters affect performance

**User Stories:**

#### 1. US-037: Run Sensitivity Analysis on Event Rate (Lambda)

- As an analyst, I want to understand load sensitivity so that performance is characterized
- Acceptance Criteria:
  - Run simulations with  $\lambda = 0.01, 0.05, 0.1, 0.15$  events/tick
  - For each  $\lambda$ : test both baseline and AI fleets
  - Collect ART and total events processed
  - Plot ART vs.  $\lambda$  for both strategies
- Estimated Points: 4
- Tasks: Implement parameter sweep, run experiments, plot results

#### 2. US-038: Run Sensitivity Analysis on Number of Ambulances

- As a fleet manager, I want to understand scaling so that fleet sizing is informed
- Acceptance Criteria:
  - Run simulations with 3, 5, 7, 10 ambulances
  - For each count: test both baseline and AI fleets
  - Collect ART and coverage metrics
  - Plot ART vs. fleet size

- Estimated Points: 4
- Tasks: Implement fleet size variation, run experiments, analyze scalability

### 3. US-039: Analyze K-Means Sensitivity (k parameter, update frequency)

- As a clustering engineer, I want optimal hotspot parameters so that dynamic dispatch is tuned
- Acceptance Criteria:
  - Test k=2, 3, 4 for clustering
  - Test update frequencies: every 25, 50, 100 ticks
  - Measure ART and rebalancing count for each combination
  - Recommend optimal configuration
- Estimated Points: 4
- Tasks: Parameter sweep for clustering, run experiments, recommend settings

### 4. US-040: Generate Sensitivity Analysis Report

- As a researcher, I want comprehensive analysis so that design choices are justified
- Acceptance Criteria:
  - Report: all sensitivity analyses with visualizations
  - Table: recommended parameters with justification
  - Trade-offs discussed (e.g., responsiveness vs. churn)
  - Future optimization directions noted
- Estimated Points: 3
- Tasks: Compile all analyses, create visualizations, write report

## Sprint 11: Integration Testing & Edge Cases (Week 11)

**Sprint Goal:** Harden system against edge cases and verify robustness

**User Stories:**

### 1. US-041: Implement Comprehensive Unit Tests

- As a QA engineer, I want unit tests so that components are reliable
- Acceptance Criteria:
  - Tests for A\*: shortest paths, unreachable nodes, single-node
  - Tests for GA: fitness calculation, mutation, convergence
  - Tests for K-Means: convergence, single-cluster, divergent data
  - Tests for Dispatcher: assignment logic, re-routing

- Minimum 80% code coverage
- Estimated Points: 6
- Tasks: Write unit tests, run coverage analysis, fix coverage gaps

## 2. US-042: Implement Integration Tests

- As a tester, I want end-to-end tests so that system behavior is correct
- Acceptance Criteria:
  - Test 1: Full simulation run (setup → 100 ticks → shutdown)
  - Test 2: GA→Dispatcher→Simulation pipeline
  - Test 3: Dynamic rebalancing triggers and works
  - Test 4: Traffic congestion causes re-routing
  - All tests pass without errors
- Estimated Points: 5
- Tasks: Design integration test scenarios, implement, run tests

## 3. US-043: Test Edge Cases & Failure Modes

- As a reliability engineer, I want edge cases handled so that system is robust
- Acceptance Criteria:
  - Empty graph / no nodes: handled gracefully
  - No idle ambulances: event queued for later
  - All ambulances busy: response time degraded but valid
  - Unreachable accident location: logged and skipped
  - Traffic congestion 100%: routing adapts
- Estimated Points: 4
- Tasks: Identify edge cases, implement handling, test each case

## 4. US-044: Create Regression Test Suite

- As a developer, I want regression tests so that refactoring is safe
- Acceptance Criteria:
  - Regression suite: 10+ test scenarios with known expected outputs
  - All baseline results recorded
  - Test suite runs in < 2 minutes
  - CI-ready (all tests pass before merge)
- Estimated Points: 3
- Tasks: Create test scenarios, record baselines, set up test runner

---

## Sprint 12: Performance Optimization & Scaling (Week 12)

**Sprint Goal:** Optimize for larger simulations and longer runs

### User Stories:

#### 1. US-045: Profile Code & Identify Bottlenecks

- As a performance engineer, I want profiling data so that optimization is targeted
- Acceptance Criteria:
  - Python profiler (cProfile) run on full simulation
  - Top 5 slowest functions identified
  - Call graph generated
  - Bottleneck root causes documented
- Estimated Points: 3
- Tasks: Run profiler, analyze output, document findings

#### 2. US-046: Optimize A Search Performance\*

- As a pathfinding engineer, I want faster A\* so that real-time routing is feasible
- Acceptance Criteria:
  - Current A\*: target < 50ms (baseline from Week 2)
  - Optimizations: better heuristic, pruning, caching
  - Benchmark: 10-path average time
  - No correctness regression
- Estimated Points: 5
- Tasks: Implement optimizations, benchmark, test correctness

#### 3. US-047: Optimize K-Means Convergence

- As a clustering engineer, I want faster convergence so that rebalancing is quicker
- Acceptance Criteria:
  - Current K-Means: target < 100ms for 100 points
  - Optimizations: early stopping, vectorization (numpy)
  - Benchmark: convergence time per run
  - Cluster quality maintained
- Estimated Points: 4

- Tasks: Vectorize code, implement early stopping, benchmark

#### 4. **US-048: Enable Simulation Scaling (10000+ ticks, 10+ ambulances)**

- As a researcher, I want long-running simulations so that patterns emerge
  - Acceptance Criteria:
    - Simulate 10000+ ticks without memory issues
    - 10+ ambulances without slowdown
    - Logging streamlined (write to disk, not memory)
    - Simulation runs in < 10 minutes
  - Estimated Points: 4
  - Tasks: Implement disk logging, optimize data structures, run long simulations
- 

## **PHASE 4: THE POLISH (Weeks 13-14)**

*Goal: Documentation, reporting, and final refinement*

### **Sprint 13: Documentation & Code Quality (Week 13)**

**Sprint Goal:** Complete documentation and ensure code quality

#### **User Stories:**

##### **1. US-049: Complete Code Documentation**

- As a maintainer, I want clear documentation so that code is understandable
- Acceptance Criteria:
  - All functions have docstrings (numpy format)
  - Module-level documentation for each .py file
  - Algorithm explanations (pseudocode for complex algorithms)
  - Examples provided for public APIs
- Estimated Points: 4
- Tasks: Write docstrings, create module READMEs, add examples

##### **2. US-050: Create Architecture Documentation**

- As an onboarding, I want system architecture so that I understand the design
- Acceptance Criteria:
  - Architecture diagram (components, data flow)
  - Design patterns documented (Entity-Component, simulation loop)

- Class hierarchy documented
- Data structure descriptions (CityGraph, DistanceMatrix, etc.)
- Estimated Points: 4
- Tasks: Create diagrams, write architecture doc, explain design choices

### 3. **US-051: Refactor Code for Readability & Maintainability**

- As a code reviewer, I want clean code so that future development is easier
- Acceptance Criteria:
  - Code style consistent (PEP 8)
  - Variable names clear and descriptive
  - Functions single-responsibility
  - No dead code or commented-out blocks
- Estimated Points: 4
- Tasks: Run linter (pylint/flake8), refactor, remove dead code

### 4. **US-052: Create User Guide & Usage Documentation**

- As an end user, I want clear instructions so that I can run the system
  - Acceptance Criteria:
    - Setup instructions (step-by-step)
    - Configuration guide (parameters, config file)
    - Running the simulation (baseline, AI, experiments)
    - Interpreting results and visualizations
  - Estimated Points: 3
  - Tasks: Write user guide, create quick-start, provide example commands
- 

## **Sprint 14: Final Report & Presentation (Week 14)**

**Sprint Goal:** Complete final project deliverables

**User Stories:**

### 1. **US-053: Compile Final Project Report**

- As an author, I want a comprehensive report so that the project is fully documented
- Acceptance Criteria:
  - Title page with project info
  - Executive summary (1 page)

- Introduction & motivation
- Technical approach (algorithms, architecture)
- Results: baseline, AI, comparisons, sensitivity analysis
- Conclusions & recommendations
- All visualizations embedded
- References and bibliography
- Estimated Points: 5
- Tasks: Write sections, compile visualizations, format report, proofread

## 2. US-054: Create Presentation Slides

- As a presenter, I want compelling slides so that the project shines
- Acceptance Criteria:
  - Title slide
  - Problem statement (motivation)
  - Technical overview (brief)
  - Key algorithms (visuals)
  - Results comparison (impact)
  - Demo / live simulation walkthrough
  - Conclusions & future work
  - 20-30 minute duration
- Estimated Points: 4
- Tasks: Design slides, gather visuals, practice presentation

## 3. US-055: Prepare Live Demo & Walkthroughs

- As a demonstrator, I want working demo so that the system is tangible
- Acceptance Criteria:
  - Quick-start script runs in < 5 minutes
  - Visualization shows ambulance movements, events, hotspots clearly
  - Optional: interactive mode to pause/resume/adjust parameters
  - Fallback: pre-recorded video if live demo fails
- Estimated Points: 3
- Tasks: Create demo script, test reproducibility, record video backup

## 4. US-056: Final Review, Polish, & Submission

- As a project lead, I want quality assurance so that deliverables are professional

- Acceptance Criteria:
    - Code review: all PRs merged, tests passing
    - Report review: spelling, grammar, formatting
    - Presentation review: clarity, flow, visuals
    - Final README with all deliverables listed
    - All files organized in submission folder
  - Estimated Points: 3
  - Tasks: Final review, address feedback, prepare submission package
- 

## Summary Table

Sprint	Week	Phase	Focus	Key Deliverables
1	1	Skeleton	Setup & Map	Environment, modeltown.graphml, distance matrix
2	2	Skeleton	Navigation	A* algorithm, test cases, path visualization
3	3	Skeleton	Optimization	GA solver, optimal stations, coverage visualization
4	4	Brain	Agents & Simulation	Ambulance class, event spawner, main loop, live viz
5	5	Brain	Dispatch	Dispatcher class, task assignment, response tracking
6	6	Brain	Clustering	K-Means implementation, hotspot detection, visualization
7	7	Brain	Realism	Traffic dynamics, re-routing, config system, logging
8	8	Pulse	Baseline	Random fleet baseline, ART measurement, comparison setup
9	9	Pulse	Optimization	AI fleet run, head-to-head comparison, results viz
10	10	Pulse	Sensitivity	Parameter analysis ( $\lambda$ , fleet size, k), report
11	11	Pulse	Testing	Unit tests, integration tests, edge cases, regression suite
12	12	Pulse	Scaling	Performance profiling, optimization, long-run capability
13	13	Polish	Documentation	Code docs, architecture docs, refactoring, user guide
14	14	Polish	Delivery	Final report, presentation, demo, submission

## Effort Distribution

- **Total Estimated Points:** 168 points (averaging 12 points per sprint)
  - **Phase 1 (Weeks 1-3):** 47 points — Foundation building
  - **Phase 2 (Weeks 4-7):** 48 points — Feature development
  - **Phase 3 (Weeks 8-12):** 56 points — Validation & optimization
  - **Phase 4 (Weeks 13-14):** 17 points — Documentation & delivery
- 

## Weekly Time Commitment

Given your 6-course load, estimated time per sprint:

- **Light Sprints (1-2, 8, 13-14):** 8-10 hours/week
- **Medium Sprints (3-5, 9-12):** 12-15 hours/week
- **Heavy Sprints (6-7, 10-11):** 15-18 hours/week

**Average:** ~12 hours/week for project work

---

## Sprint Planning Tips

1. **Backlog Refinement:** Review user stories before each sprint
  2. **Estimation:** Use planning poker or T-shirt sizing for better accuracy
  3. **Standups:** Quick daily 15-min check-ins (even solo, helps track progress)
  4. **Retrospective:** End-of-sprint reflection on what went well, what to improve
  5. **Buffer:** Allocate 10-20% contingency for unexpected issues
- 

## Success Criteria

- ✓ All user stories completed by end of sprint
- ✓ Code tested (unit + integration)
- ✓ Documentation updated
- ✓ No technical debt carried forward
- ✓ Final project demonstrates 20%+ ART improvement over baseline